



Chaos Mesh®



Chaos Mesh 开源说

云原生混沌工程平台



关于我们

姓名：杨可奥

GitHub 地址

: <https://github.com/YangKeao/>

公司、职位：PingCAP 研发



关于我们


姓名: 周强

GitHub 地址


: <https://github.com/zhouqiang-cl>

公司、职位: PingCAP 工程效率负责人,
ChaosMesh 负责人






Part I - 混沌工程和 Chaos Mesh



分布式系统越来越复杂



amazon.com®




NETFLIX



Chaos Mesh®

chaos-mesh.org

故障随时都有可能发生

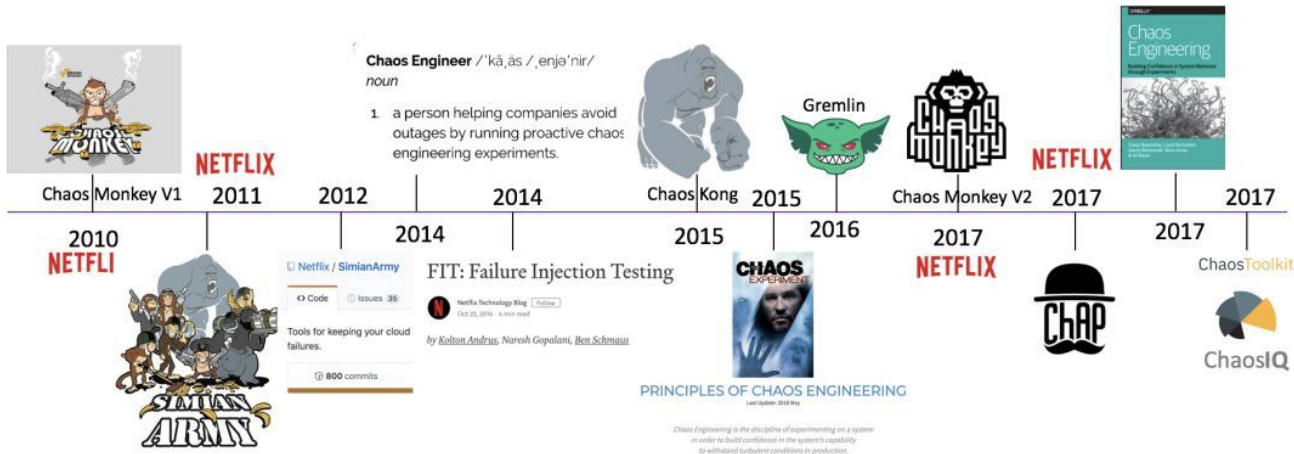


Chaos Mesh®

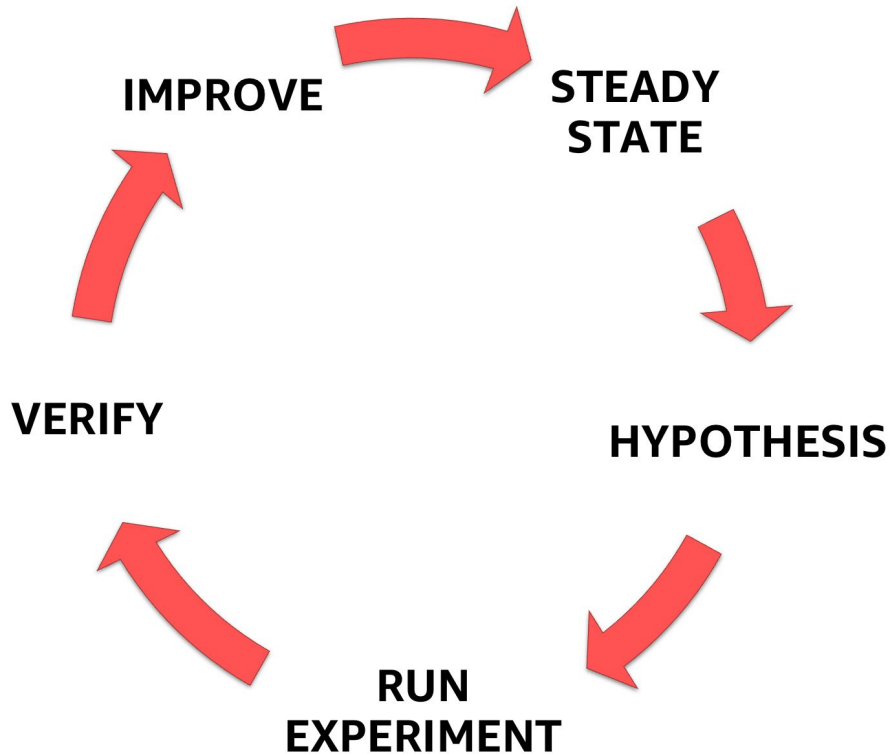
chaos-mesh.org

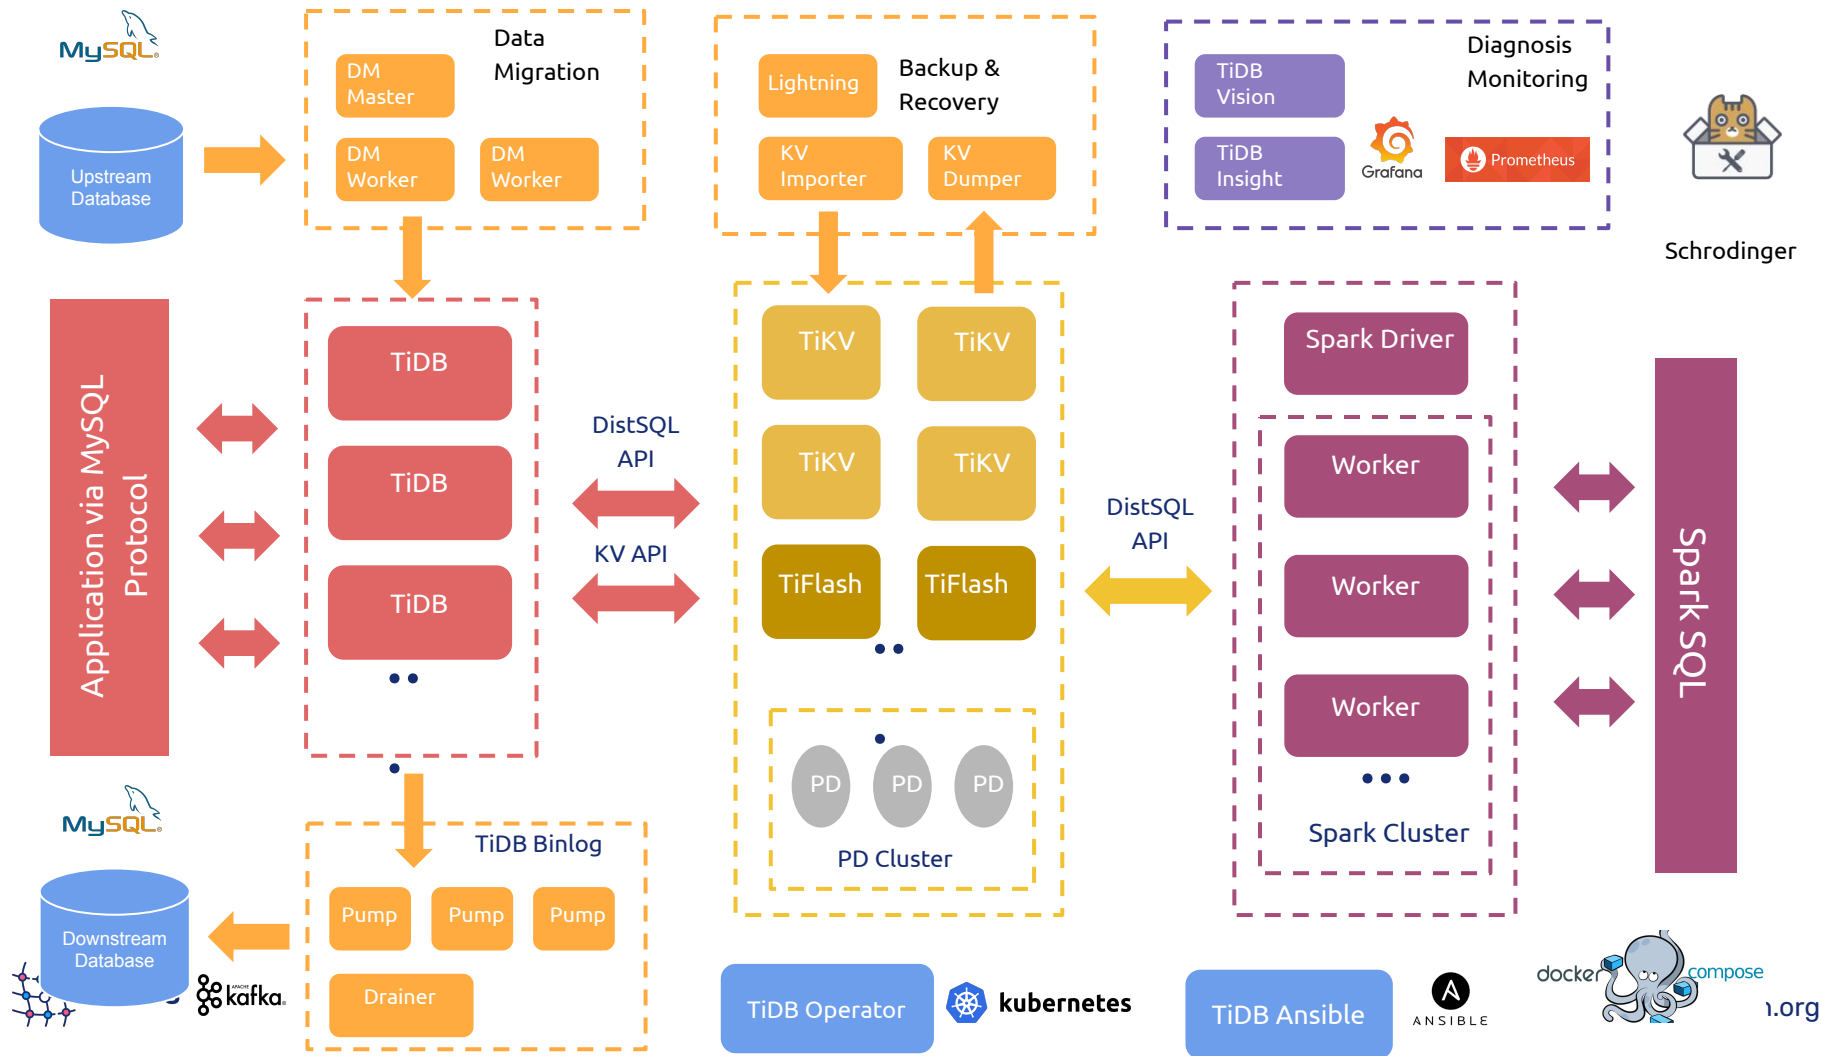
混沌工程定义

混沌工程是一门新兴的技术学科，他的初衷是通过实验性的方法，让人们建立对于复杂分布式系统在生产中抵御突发事件能力的信心。



混沌工程步骤







Chaos Monitor

chaos events



short-term
pod-failure

long-term
pod-failure


Duration

unexpected high duration





Here we have 2 short-time and a long-time pod-failures of TiKV pod

混沌工程在 PingCAP 的实践



Schrodinger's Cat



Chaos Mesh®



Chaos Mesh®

Community



3,400+

Stars




86+

Contributors
and growing




Sandbox

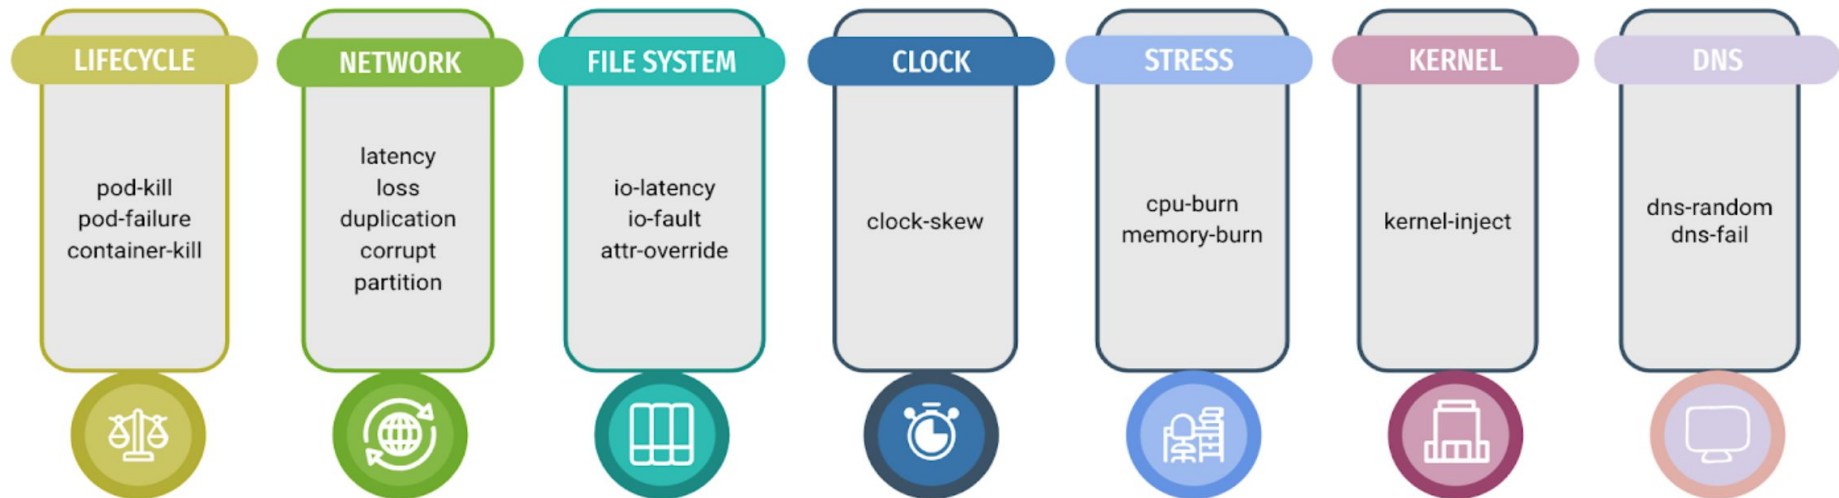
Cloud Native
Computing Foundation




Part II - Chaos Mesh实现原理



Chaos Mesh 功能



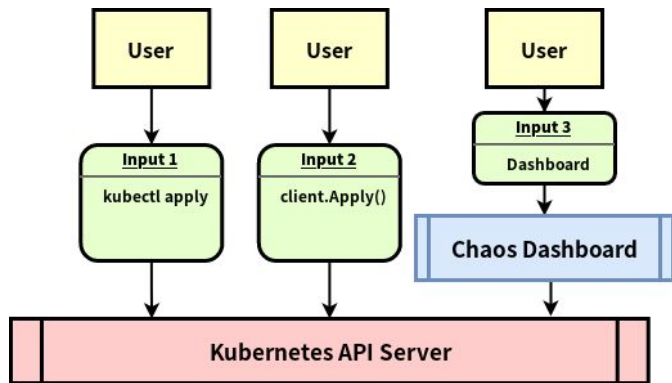
Chaos Mesh 整体架构



- 用户输入、观测
- 监听资源变化, 进行注入/恢复
- 在具体节点上进行故障注入



用户输入、观测



- 使用 kubectl 工具提交
- 使用 Kubernetes Client
- 使用 Chaos Dashboard

使用 kubectl 工具提交

```
apiVersion: chaos-mesh.org/v1alpha1
kind: NetworkChaos
metadata:
  name: network
  namespace: chaos-testing
spec:
  action: partition
  mode: one
  selector:
    labelSelectors:
      "app.kubernetes.io/component": "tikv"
  direction: to
  target:
    selector:
      labelSelectors:
        "app.kubernetes.io/component": "tikv"
    mode: one
  duration: "10s"
  scheduler:
    cron: "@every 15s"
```

- `kubectl apply -f ./network.yaml`
- `kubectl describe NetworkChaos network`

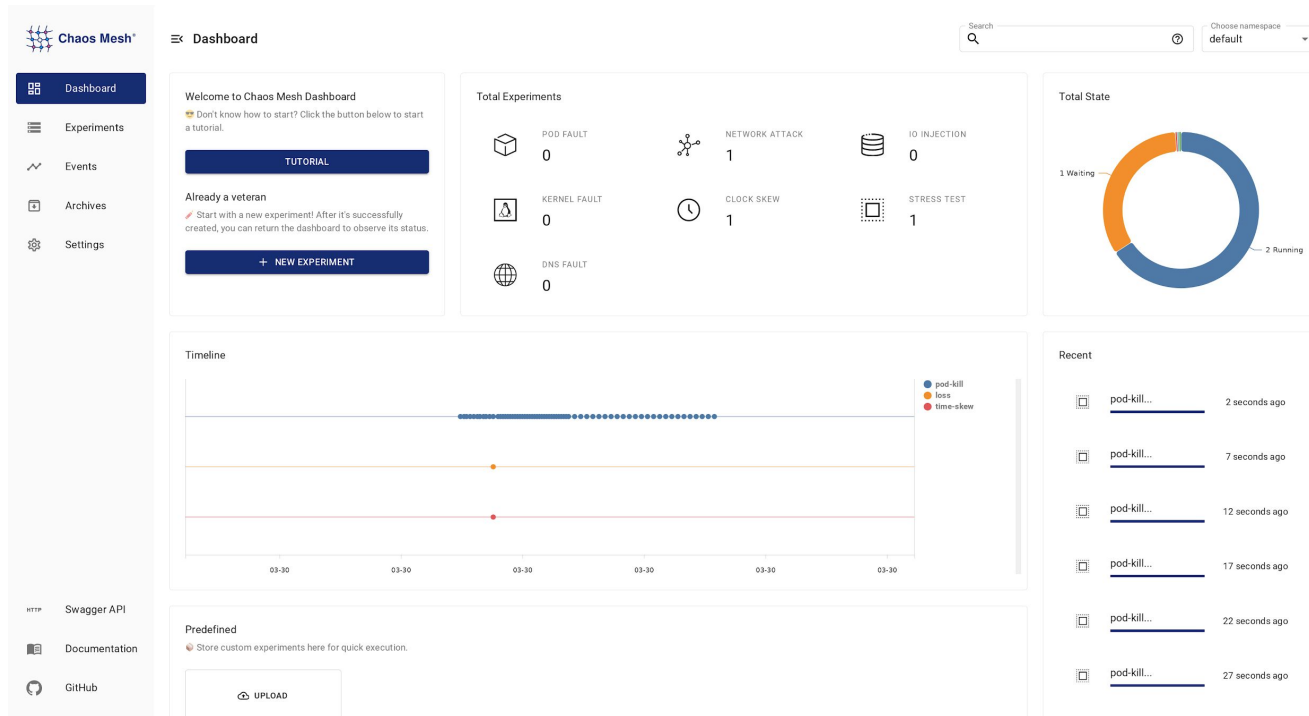


使用 Kubernetes API/Client


```
podFailureChaos := &v1alpha1.PodChaos{
  ObjectMeta: metav1.ObjectMeta{
    Name:      "timer-failure",
    Namespace: ns,
  },
  Spec: v1alpha1.PodChaosSpec{
    Selector: v1alpha1.SelectorSpec{
      Namespaces: []string{
        ns,
      },
      LabelSelectors: map[string]string{
        "app": "timer",
      },
    },
    Action: v1alpha1.PodFailureAction,
    Mode:   v1alpha1.OnePodMode,
  },
}
err = cli.Create(ctx, podFailureChaos)
```

- 方便集成入测试流程
- 可编程地动态控制注入流程

Chaos Dashboard

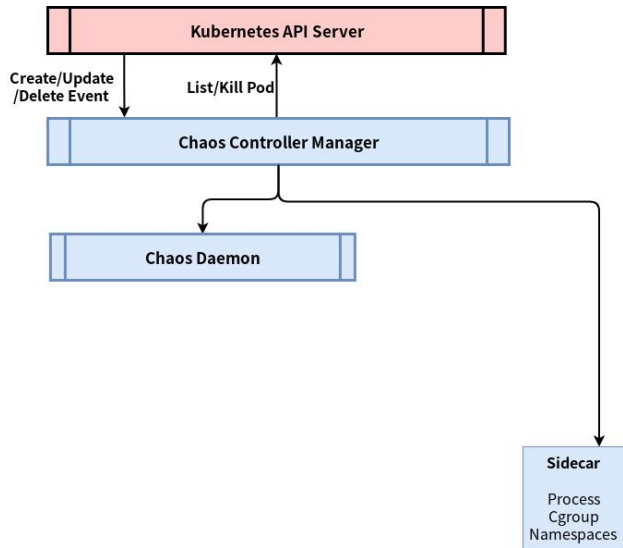


Chaos Dashboard



- 友好的用户界面
- 接入 RBAC 的权限管控
- 方便管理和观察已有的错误

监听资源的变化



- 监听 PodChaos, NetworkChaos... 等资源的 创建/更新/删除
- 决定当前该 注入 / 恢复 / 等待
- (进行简单的注入, 比如 PodKill)
- 向 Chaos Daemon / Sidecar 发送请求

Reconcile



Resource		Physical Implementation
ReplicaSet	↔	Pod
Pod	↔	Container
ConfigMap	↔	key/value in etcd
...		...
NetworkChaos	✗ →	iptables rules / tc qdisc / ...


Reconcile 以 ReplicaSet 为例

Name: hello-kubernetes-74c4d446d7
Namespace: default
Controlled By: Deployment/hello-kubernetes
Replicas: 0 current / 3 desired
Pods Status: 0 Running / 0 Waiting / 0 Succeeded / 0
Failed
Pod Template:
Labels: app=hello-kubernetes
pod-template-hash=74c4d446d7
Containers:
hello-kubernetes:
Image: paulbouwer/hello-kubernetes:1.8
Port: 8080/TCP
Host Port: 0/TCP
Limits:
cpu: 50m
Environment: <none>
Mounts: <none>
Volumes: <none>
Events: <none>




Name: hello-kubernetes-74c4d446d7
Namespace: default
Controlled By: Deployment/hello-kubernetes
Replicas: 3 current / 3 desired
Pods Status: 3 Running / 0 Waiting / 0 Succeeded / 0
Failed
Pod Template:
Labels: app=hello-kubernetes
pod-template-hash=74c4d446d7
Containers:
hello-kubernetes:
Image: paulbouwer/hello-kubernetes:1.8
Port: 8080/TCP
Host Port: 0/TCP
Limits:
cpu: 50m
Environment: <none>
Mounts: <none>
Volumes: <none>
Events: <none>

Reconcile 以 ReplicaSet 为例



- Pod 和 ReplicaSet 的变化都会触发一次 Reconcile
- 更新 Pod 状态
- 比较 current replicas 和 desired replicas
- 如有必要, 创建新的 Pod
- 更新 ReplicaSet 的状态

Reconcile 以 *Chaos 为例



```
apiVersion: chaos-mesh.org/v1alpha1
kind: NetworkChaos
metadata:
  name: network
  namespace: chaos-testing
spec:
  action: partition
  mode: one
  selector:
    labelSelectors:
      "app.kubernetes.io/component": "tikv"
  direction: to
  target:
    selector:
      labelSelectors:
        "app.kubernetes.io/component": "tikv"
  mode: one
  duration: "10s"
  scheduler:
    cron: "@every 15s"
```

- 是不是已经删除了？
- 是不是注入过了？
- 需不需要 cronly 多次运行？
- 该如何注入 / 恢复？

Reconcile 以 *Chaos 为例

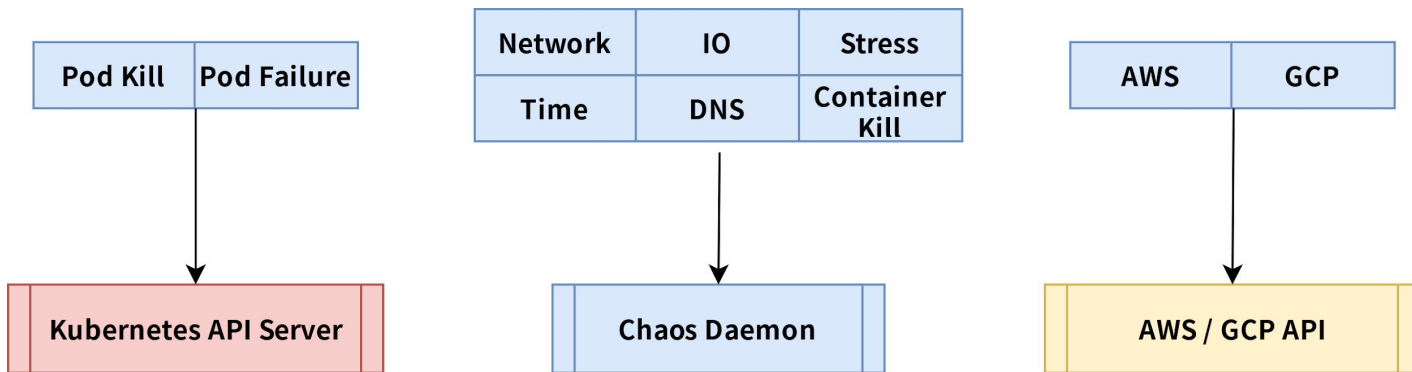
now: Wed Mar 31 02:39:38 PM CST 2021

```
duration: "10s"  
scheduler:  
  cron: "@every 15s"
```

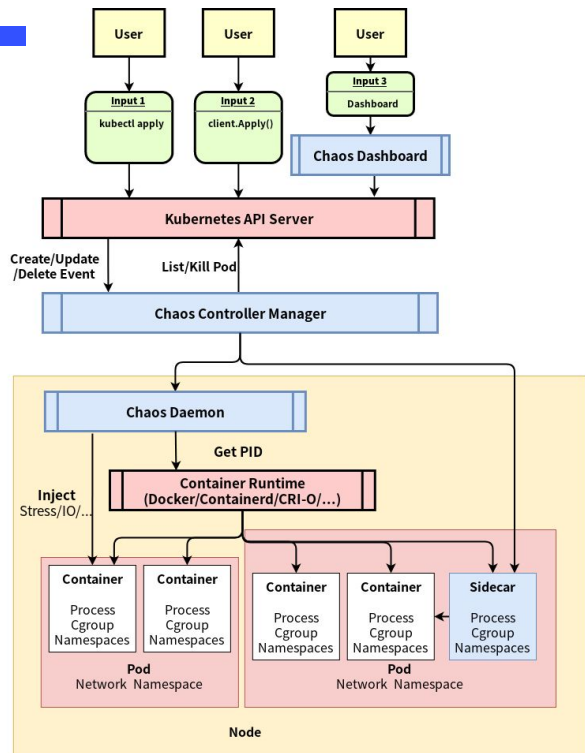
```
status:  
  nextStart: "Wed Mar 31 02:39:37 PM CST 2021"  
  nextRecover: "Wed Mar 31 02:50:37 PM CST 2021"
```

- 是不是已经删除了？
- 是不是注入过了？
- 需不需要 cronly 多次运行？

*Chaos 如何注入？



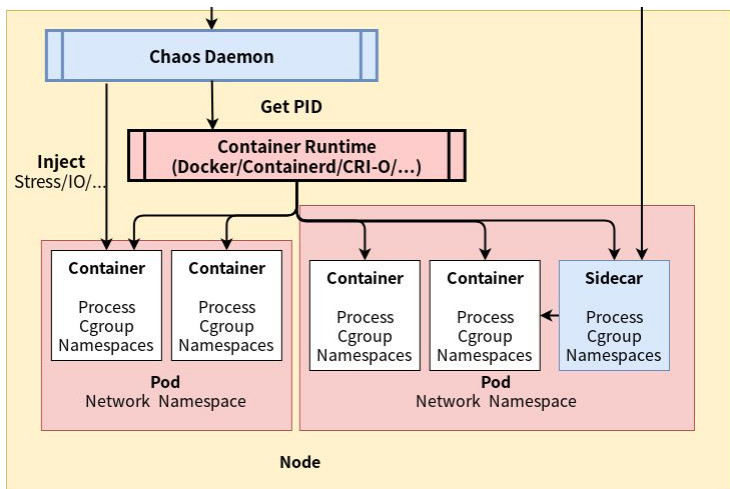
Chaos Mesh 整体架构



- 用户输入、观测
- 监听资源变化, 进行注入/恢复
- 在具体节点上进行故障注入



Chaos Daemon 注入原理




- Container 的实体：
 - 进程
 - Namespace
 - Cgroup

控制可见性
限制资源


} 隔离
- 注入的实质
 - 侵入 Namespace / Cgroup
 - 进行干扰、注入

Cgroup 影响资源分配

- CPU 限制 CPU 资源分配
- Memory 限制 内存 资源分配
- Pid 限制 Pid 数量资源的分配
- ...



侵入 namespace



Namespace 影响查询过程

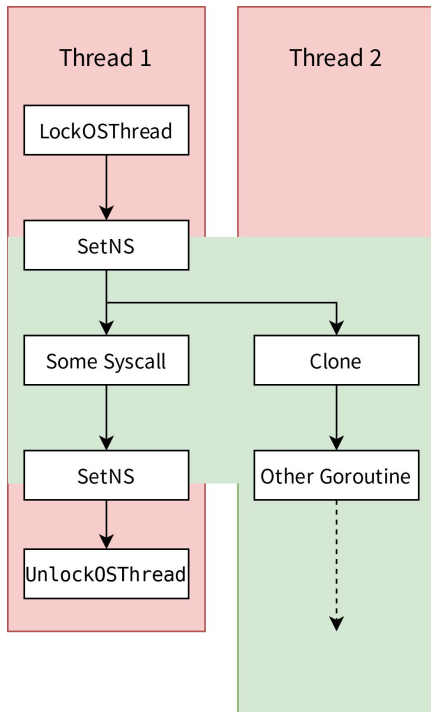
- mnt namespace 影响 path resolution
- pid namespace 影响 process lookup
- net namespace 影响 network device lookup
- ...

Linux namespace and Go Don't Mix

```
1: [pid 3361] openat(AT_FDCWD, "/proc/17526/ns/net", O_RDONLY) = 61
2: [pid 3361] getpid() = 3357
3: [pid 3361] gettid() = 3361
4: [pid 3361] openat(AT_FDCWD, "/proc/3357/task/3361/ns/net", O_RDONLY) = 62
5: [pid 3361] setns(61, CLONE_NEWNET) = 0
<...>
6: [pid 3361] socket(AF_NETLINK, SOCK_RAW, NETLINK_ROUTE) = 63
7: [pid 3361] bind(63, {sa_family=AF_NETLINK, pid=0, groups=00000000}, 12) = 0
8: [pid 3361] sendto(63, "\x20\x00...", 32, 0, {sa_family=AF_NETLINK, pid=0, groups=00000000}, 12) = 32
9: [pid 3361] getsockname(63, {sa_family=AF_NETLINK, pid=3357, groups=00000000}, [12]) = 0
10: [pid 3361] futex(0xc820504110, FUTEX_WAKE, 1 <unfinished ...>
11: [pid 3361] <... futex resumed> ) = 1
12: [pid 3361] futex(0xd82930, FUTEX_WAKE, 1) = 1
13: [pid 3361] futex(0xc820060110, FUTEX_WAIT, 0, NULL <unfinished ...>
14: [pid 3361] <... futex resumed> ) = 0
15: [pid 3361] recvfrom(63, <unfinished ...>
16: [pid 3361] <... recvfrom resumed> "\x4c\x00...", 4096, 0, {sa_family=AF_NETLINK, pid=0, groups=00000000}, [12]) = 236
<...>
17: [pid 3361] clone( <unfinished ...>
18: [pid 3361] <... clone resumed> child_stack=0x7f19efffee70,
flags=CLONE_VM|CLONE_FS|CLONE_FILES|CLONE_SIGHAND|CLONE_THREAD|CLONE_SYSVSEM|CLONE_SETTLS|CLONE_PARENT
T_SETTID|CLONE_CHILD_CLEARPID, parent_tidptr=0x7f19effff9d0, tls=0x7f19effff700, child_tidptr=0x7f19effff9d0) = 3365
<...>
19: [pid 3361] setns(62, CLONE_NEWNET <unfinished ...>
20: [pid 3361] <... setns resumed> ) = 0
<...>
21: [pid 3365] sendto(65, "\x2c\x00...", 44, 0, {sa_family=AF_NETLINK, pid=0, groups=00000000}, 12) = 44
22: [pid 3365] getsockname(65, {sa_family=AF_NETLINK, pid=3357, groups=00000000}, [12]) = 0
23: [pid 3365] recvfrom(65, "\x40\x00...", 4096, 0, {sa_family=AF_NETLINK, pid=0, groups=00000000}, [12]) = 64
24: [pid 3365] close(65) = 0
25: [pid 3365] write(2, "Cannot find weave bridge: Link not found\n", 41
```



Linux namespace and Go Don't Mix

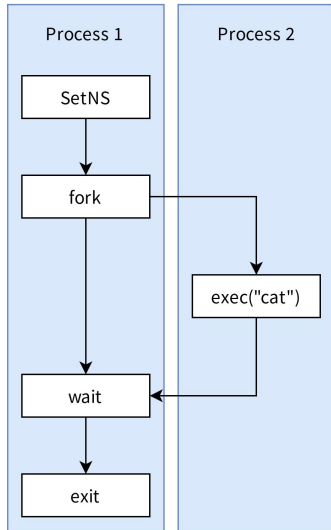


Namespace 泄漏


会导致其他线程的 Namespace 难以追踪

nsenter

Pid Namespaces



Other Namespaces



`nsenter --pid /proc/xxx/ns/pid cat`

nsenter 的问题

1. 信号处理较为随意, 难以跟踪子进程

2. 难以应对 mnt namespace 的情形:

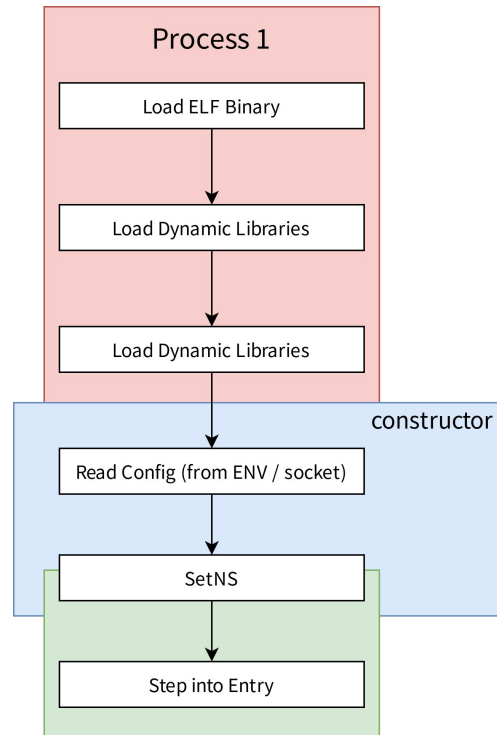
```
nsenter --mnt /proc/xxx/ns/mnt cat
```

```
nsenter: failed to execute cat: No such file or directory
```

这一情形在 Distroless 容器逐渐流行的今天更加普遍了

runc 方案

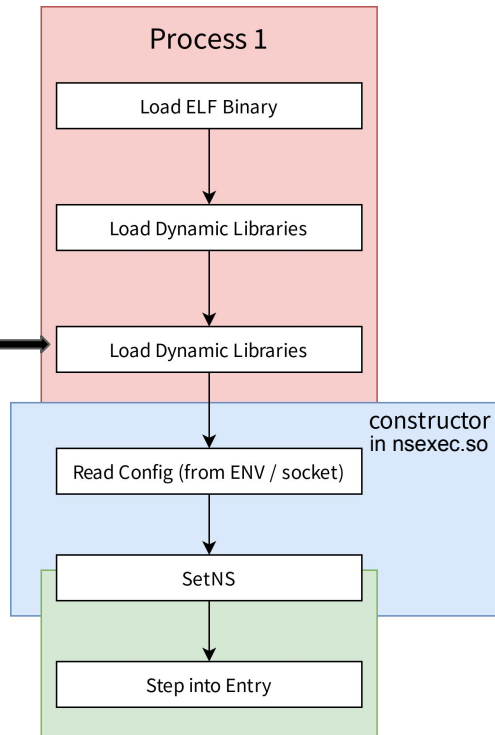
```
#cgo CFLAGS: -Wall
extern void nsexec();
void __attribute__((constructor)) init(void) {
    setns();
}
```



chaos-mesh/nsexec 方案


1. 设置 LD_PRELOAD
2. 在链接库的 constructor 中执行 setns

LD_PRELOAD=nsexec.so




侵入 namespace

- `nsexec --net=/proc/xxx/ns/net iptables -A`
- `nsexec --net=/proc/xxx/ns/net tc qdisc add`
- `nsexec --pid=/proc/xxx/ns/pid stress-ng ...`
- `nsexec --mnt=/proc/xxx/ns/mnt inject ...`




从流程中注入




以(入方向)限流为例:tc qdisc

First-in First-out (FIFO)



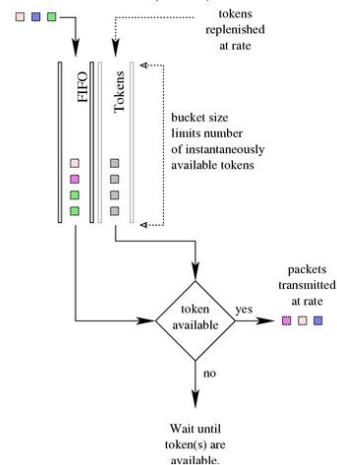
FIFO

pfifo_fast queuing discipline



pfifo_fast / prio

Token Bucket Filter (TBF)



tbf



以时间偏移为例

- ``clock_gettime`` syscall (228 on x86_64)
- ``clock_gettime`` vDSO call
- Language/Runtime specific function call ★
 - Rust: ``Instant::now()``
 - Go: ``time.Now()``
 - C: ``clock_gettime`` in glibc
 - ...

language/runtime 专门的函数调用

- C: ``clock_gettime`` in glibc → ``clock_gettime`` vDSO call
- Rust: ``Instant::now()`` → ``clock_gettime`` in glibc → ``clock_gettime`` vDSO call
- Go: ``time.Now()`` → ``clock_gettime`` vDSO call
-

``clock_gettime`` vDSO call 覆盖了大部分情况！


什么是 vDSO ?

- 由 kernel 提供在用户态运行的函数的机制
 - `clock_gettime`
 - `getcpu`
 - `gettimeofday`
 - `time`
- 一段由 kernel 在程序启动时分配的内存区域
- 这部分内存的内容和动态链接库的 ELF 格式一致

如何完成注入？

1. `PTTRACE_ATTACH` 容器中的每一个进程
2. `PTTRACE_POKEDATA` 修改 vDSO 中时间相关函数的实现
3. `PTTRACE_DETACH`

目标进程
内存空间






thix

浙江 杭州



扫一扫上面的二维码图案，加我微信

Join us !!!



Chaos Mesh 群组

